



# Training a Neural Network to Predict Stock Prices

**BY: YOUSSEF ELMOUGY**





1


# PROBLEM DESCRIPTION

Utilize a huge dataset containing stock prices for the S&P 500 Index and its constituents to precisely predict the next minute stock price of the S&P 500 Index

A decorative graphic on the left side of the slide. It features a large cyan hexagon in the center containing the number '2'. Surrounding this central hexagon are several smaller hexagons of varying shades of blue and cyan. Some of these smaller hexagons contain white icons: a lightbulb, a thumbs-up, a smartphone, a magnifying glass, and a gear. There is also a network-like icon with a central node and radiating lines. The entire graphic is set against a dark blue background.


2

# DATA DESCRIPTION

- 
- Dataset contains 41,266 minutes of data ranging from April 2017 to August 2017 on prices of the 500 stock constituents along with the total S&P 500 index price

	SP500	NASDAQ.AAL	...	NYSE.ZBH	NYSE.ZTS
count	41266.000000	41266.000000	...	41266.000000	41266.000000
mean	2421.537882	47.708346	...	121.423515	60.183874
std	39.557135	3.259377	...	5.607070	3.346887
min	2329.139900	40.830000	...	110.120000	52.300000
25%	2390.860100	44.945400	...	117.580000	59.620000
50%	2430.149900	48.360000	...	120.650000	61.585600
75%	2448.820100	50.180000	...	126.000000	62.540000
max	2490.649900	54.475000	...	133.450000	63.840000

Fig. 1: `print(dataset.describe())`

- Each row of the dataset contains the constituent 500 stock's prices at time  $T = t$  and the stock price of the S&P 500 at  $T = t + 1$
- 

A decorative pattern of hexagons in various shades of blue and cyan. Some hexagons contain white icons: a lightbulb, a thumbs-up, a smartphone, a magnifying glass, and a gear. A network of dots is also visible on the left side.

3


# METHOD DESCRIPTION

This deep learning model is built with TensorFlow



The dataset is split into 80% used as the training data, and 20% used as the testing data. Both the training and testing data are scaled using sklearn's **MinMaxScaler()** and bounded within the range  $[-1, 1]$

The TensorFlow model consists of four layers. The number of neurons in each layer is adjusted to find the architecture that produces the best accuracy.



Each subsequent layer's number of neurons is always half the number of neurons of the previous layer



Information is compressed as it flows between layers hence creating a more reliable accuracy






The model is represented through placeholders and variables:

◇ placeholders


- **ins** : contains the NN's inputs (the stock prices of the 500 constituents)
- **outs** : contains the NN's outputs (the stock price of the S&P 500)



◇ variables – each layer, including the output layer, has a unique set of variables:

- Weight variable : each layer passes its output as the input of the next layer
- Bias variable : the number of neurons in the layer

The placeholders and variables are then combined to design the architecture of the neural network





The NN is then fitted and trained using adjustable sized batches

For each batch,

- The error is calculated using the Mean Squared Error (MSE) approach
- An optimizer is applied to minimize the MSE
- The predictions of the NN are plotted against the actual stock prices





# 4

## Implementation

Details on how the algorithm was implemented

Importing the data

Splitting the dataset into training data and testing data

Scaling the dataset

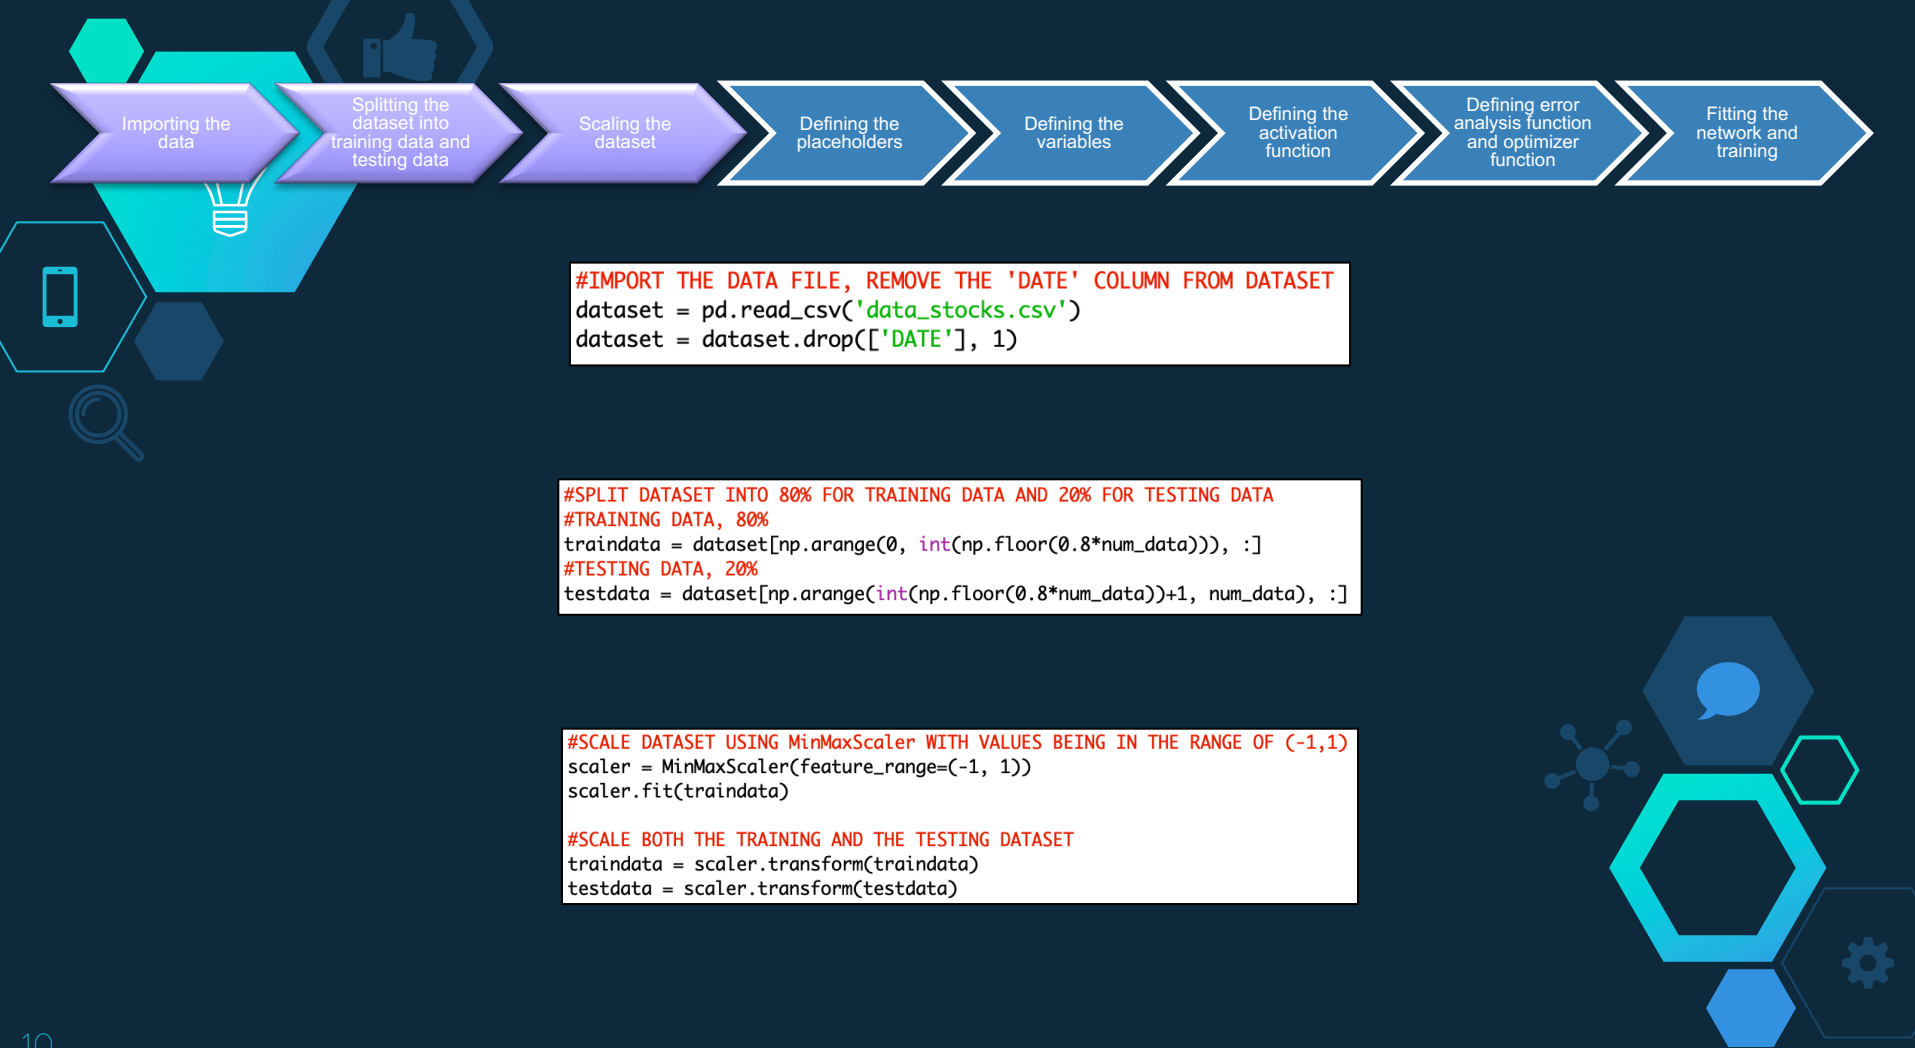
Defining the placeholders

Defining the variables

Defining the activation function

Defining error analysis function and optimizer function

Fitting the network and training





**# placeholders**

```
ins = tf.placeholder(dtype=tf.float32, shape=[None, num_stocks])  
outs = tf.placeholder(dtype=tf.float32, shape=[None])
```





```
# layeri_neurons, -----TRY OUT DIFFERENT NUMBER OF NEURONS-----  
layer1_neurons = 1000 # double input size  
layer2_neurons = 500 # 50% of previous layer  
layer3_neurons = 250 # 50% of previous layer  
layer4_neurons = 125 # 50% of previous layer
```

```
# layeri_weight, layeri_bias  
layer1_weight = tf.Variable(weight_initializer([num_stocks, layer1_neurons]))  
layer1_bias = tf.Variable(bias_initializer([layer1_neurons]))  
layer2_weight = tf.Variable(weight_initializer([layer1_neurons, layer2_neurons]))  
layer2_bias = tf.Variable(bias_initializer([layer2_neurons]))  
layer3_weight = tf.Variable(weight_initializer([layer2_neurons, layer3_neurons]))  
layer3_bias = tf.Variable(bias_initializer([layer3_neurons]))  
layer4_weight = tf.Variable(weight_initializer([layer3_neurons, layer4_neurons]))  
layer4_bias = tf.Variable(bias_initializer([layer4_neurons]))  
output_weight = tf.Variable(weight_initializer([layer4_neurons, 1]))  
output_bias = tf.Variable(bias_initializer([1]))
```





```
layer1 = tf.nn.relu(tf.add(tf.matmul(ins, layer1_weight), layer1_bias))
layer2 = tf.nn.relu(tf.add(tf.matmul(layer1, layer2_weight), layer2_bias))
layer3 = tf.nn.relu(tf.add(tf.matmul(layer2, layer3_weight), layer3_bias))
layer4 = tf.nn.relu(tf.add(tf.matmul(layer3, layer4_weight), layer4_bias))
layer_output = tf.transpose(tf.add(tf.matmul(layer4, output_weight), output_bias))
```

```
layer1 = tf.nn.tanh(tf.add(tf.matmul(ins, layer1_weight), layer1_bias))
layer2 = tf.nn.tanh(tf.add(tf.matmul(layer1, layer2_weight), layer2_bias))
layer3 = tf.nn.tanh(tf.add(tf.matmul(layer2, layer3_weight), layer3_bias))
layer4 = tf.nn.tanh(tf.add(tf.matmul(layer3, layer4_weight), layer4_bias))
layer_output = tf.transpose(tf.add(tf.matmul(layer4, output_weight), output_bias))
```

```
layer1 = tf.nn.sigmoid(tf.add(tf.matmul(ins, layer1_weight), layer1_bias))
layer2 = tf.nn.sigmoid(tf.add(tf.matmul(layer1, layer2_weight), layer2_bias))
layer3 = tf.nn.sigmoid(tf.add(tf.matmul(layer2, layer3_weight), layer3_bias))
layer4 = tf.nn.sigmoid(tf.add(tf.matmul(layer3, layer4_weight), layer4_bias))
layer_output = tf.transpose(tf.add(tf.matmul(layer4, output_weight), output_bias))
```





```
#ERROR ANALYSIS FUNCTION, Measure of deviation of predictions and actual using Mean Squared Error
MSE = tf.reduce_mean(tf.squared_difference(layer_output, outs))
trainMSE = []
testMSE = []
#OPTIMISER RATE TO DECREASE THE MSE, using Adaptive Moment Estimation Optimizer (default for deep learning dev)
MSE_dec = tf.train.AdamOptimizer().minimize(MSE)
```





Importing the data

Splitting the dataset into training data and testing data

Scaling the dataset

Defining the placeholders

Defining the variables


Defining the activation function

Defining error analysis function and optimizer function

Fitting the network and training

**#TRAINING WITH DIFFERENT SIZED BATCHES FOR EACH EPOCH**

```
for epoch in range(10):  
    #GENERATE SHUFFLED TRAINING DATA  
    size = len(y_train)  
    batch_range = size // 256  
    random = np.random.permutation(np.arange(size))  
    X_train = X_train[random]  
    y_train = y_train[random]  
    for x in range(0, batch_range):  
        #TRAIN AND RUN THE BATCH AND MINIMIZE MSE  
        X_batch = X_train[(256*x):(256*x)+256]  
        Y_batch = y_train[(256*x):(256*x)+256]  
        session.run(MSE_dec, feed_dict={ins:X_batch, outs:Y_batch})  
  
    #DISPLAY PLOT EVERY 50th BATCH  
    if(np.mod(x, 50) == 0):  
        #RUN A PREDICTION ON THE DATA  
        prediction = session.run(layer_output, feed_dict={ins: X_test})  
        pred_line.set_ydata(prediction)  
        plt.pause(0.01)
```



A decorative graphic on the left side of the slide. It features a large cyan hexagon in the center containing the number '5'. Surrounding this central hexagon are several smaller hexagons of varying shades of blue and cyan. Some of these smaller hexagons contain white icons: a lightbulb, a thumbs-up, a smartphone, a magnifying glass, and a gear. There is also a network-like icon with a central node and five connecting lines. The entire graphic is set against a dark blue background.

5

RESULTS



Layer1 = 1000, layer2 = 500, layer3 = 250, layer4 = 125

MSE for test data: 0.004365  
Accuracy on test data: 0.9956350000575185

Layer1 = 500, layer2 = 250, layer3 = 125, layer4 = 100

MSE for test data: 0.004756349  
Accuracy on test data: 0.9952436508610845

Layer1 = 2000, layer2 = 1000, layer3 = 500, layer4 = 250

MSE for test data: 0.0020091033  
Accuracy on test data: 0.9979908966924995

Layer1 = 2000, layer2 = 1000, layer3 = 500, layer4 = 250

MSE for test data: 0.007297084  
Accuracy on test data: 0.9927029157988727

Layer1 = 500, layer2 = 250, layer3 = 125, layer4 = 100

MSE for test data: 0.0070824847  
Accuracy on test data: 0.9929175153374672

Layer1 = 1000, layer2 = 500, layer3 = 250, layer4 = 125

MSE for test data: 0.005050706  
Accuracy on test data: 0.9949492937885225

Activation Function: ReLU

Layer1 = 2000, layer2 = 1000, layer3 = 500, layer4 = 250

MSE for test data: 0.022667188  
Accuracy on test data: 0.9773328118026257

Layer1 = 1000, layer2 = 500, layer3 = 250, layer4 = 125

MSE for test data: 0.004171451  
Accuracy on test data: 0.9958285489119589

Layer1 = 500, layer2 = 250, layer3 = 125, layer4 = 100

MSE for test data: 0.0037012252  
Accuracy on test data: 0.9962987748440355

Activation Function: sigmoid



Activation Function: tanh

A decorative graphic on the left side of the slide. It features a large cyan hexagon with the number '6' inside. Surrounding this central hexagon are several smaller hexagons of varying shades of blue and cyan. Some of these smaller hexagons contain white icons: a lightbulb, a thumbs-up, a smartphone, a magnifying glass, and a gear. There is also a network-like icon with a central node and several connecting lines.

6

# CONCLUSION

- ◇ Decreasing the number of neurons in each layer does not necessarily increase the accuracy on test data
- ◇ Effectiveness (greatest accuracy on test data):

**sigmoid**

<

**tanh**

<

**ReLU**



The architecture combination that yields to the best accuracy results is the following:


**Activation Function: ReLU**

**layer1 = 2000, layer2 = 1000, layer3 = 500, layer4 = 250**

**Accuracy on test data: 0.9979908966924995**




Figure 1



Actual Stock  
Prices

Predicted  
Stock Prices





- ◇ The NN quickly adapts and continues to find and learn finer patterns of the data
- ◇ The optimizer works to reduce the learning rate as the model trains
  - Reduces the chance of overshooting maximum accuracy
- ◇ After 10 epochs, the data was pretty much close to a perfect fit
  - **Final MSE = 0.0020091033**





Thanks!

